

Fixed-Point Representation

Fixed-point representation allows us to use fractional numbers on low-cost integer hardware.

Problem: Integer Hardware Can't Handle Fractions

Most hardware (like FPGAs or microcontrollers) natively supports **integer arithmetic** but has **limited or no support for floating-point operations** (unless specialized hardware like an FPU is present). Directly storing fractions in integer hardware isn't possible.

Solution: Multiply by a Power of 2

By multiplying a real number by a power of 2, we shift the decimal point, converting it into an integer that can be stored and manipulated efficiently.

$$\text{Fixed-Point Value} = \text{Real Value} \times 2^N$$

where N is the number of fractional bits.

Why Not Use Floating-Point?

- Floating-point requires specialized hardware (FPUs)
- It consumes more power and FPGA resources
- Fixed-point (Q format) is much faster and efficient

What is Q Format?

The Q format represents real numbers in **fixed-point notation** using a specific number of bits for the integer and fractional parts.

Notation:

$$Q_{m.n}$$

where:

- m = Number of **integer bits** (includes sign bit)
- n = Number of **fractional bits**

Texas Instruments version

The Q notation, as defined by [Texas Instruments](#),^[1] consists of the letter Q followed by a pair of numbers $m.n$, where m is the number of bits used for the integer part of the value, and n is the number of fraction bits.

By default, the notation describes *signed* binary fixed point format, with the unscaled integer being stored in [two's complement](#) format, used in most binary processors. The first bit always gives the sign of the value (1 = negative, 0 = non-negative), and it is *not* counted in the m parameter. Thus, the total number w of bits used is $1 + m + n$.

Example 1:

Assume that an algorithm tested using floating-point arithmetic involves operations on $a = 9.216957436091198_{10}$. Now that we are satisfied with the performance of the algorithm in floating-point representation, we have decided to implement it on a low-cost fixed-point processor which has a wordlength of 16 bits. What would be the appropriate Q format to represent a on this processor?

$$\min(m+1) = 5$$

$$m=4, n=11$$

Q4.11 format

the Q4.11 representation of a without the implied binary point is equal to a multiplied by 2^{11} . Hence, to represent a in the Q4.11 format, we multiply it by 2^{11} , round it to the nearest integer, and convert the rounded result into the binary form.

$$a \times 2^{11} = 18876.3288 \approx 18876 = 100\ 1001\ 1011\ 1100_{(2)}$$

Since a is positive, we only need to consider a sign bit of zero. Therefore, the $Q_{11}^{4.1}$ format of the number will be 01001.00110111100. For a negative number, we would have to first find the Q format of the absolute value and, then, convert it to the two's complement representation to take negative sign into account.

To represent **-5.354** in **Q3.4** format

Step 1: Understanding Q3.4 Format

- Q3.4 means:
 - 3 integer bits (including sign bit).
 - 4 fractional bits.
 - Total = 8 bits (1 sign bit + 2 integer bits + 4 fractional bits).
- The range of Q3.4 format:
 - Minimum value: -8 (1000.0000)
 - Maximum value: +7.9375 (0111.1111)
 - Resolution (step size): $2^{-4} = 0.0625$

Q3.4 representation of -5.354:

10101010₂

This 8-bit binary number represents -5.354 in Q3.4 format.

Step 2: Convert -5.354 to Fixed-Point

Multiply by $2^4 = 16$

$$-5.354 \times 16 = -85.664$$

Round to nearest integer:

$$-86$$

Step 3: Convert -86 to 8-bit Two's Complement

1. Convert 86 to binary (unsigned 8-bit):

$$86_{10} = 01010110_2$$

2. Find Two's Complement (Negate):

- Invert bits: 1010 1001
- Add 1: 1010 1010 (This is -86 in two's complement)

Verification

- Convert `1010 1010` back to decimal:
 1. Two's complement of `1010 1010`:
 - Invert bits: `0101 0101`
 - Add 1: `0101 0110` = 86
 2. Interpret as Q3.4:

$$86/16 = 5.375$$

3. Small rounding error: Original value was -5.354, stored as -5.375 due to rounding.

Now, convert some real numbers from the range **-5 to 5** into **Q3.4** format:

Real Value	Multiply by 16	Binary (8-bit)	Hex
-5.0	-80	1011 0000	B0
-3.5	-56	1100 1000	C8
-2.0	-32	1110 0000	E0
-1.25	-20	1110 1100	EC
0.0	0	0000 0000	00
1.5	24	0001 1000	18
3.75	60	0011 1100	3C
5.0	80	0101 0000	50

Choosing a Scaling Factor (Q-Format)

- The scaling factor is determined by how much **fractional precision** is needed.
- **Higher scaling factor** (e.g., 2^{14} for Q2.14) means **more precision** but **smaller range**.
- **Lower scaling factor** (e.g., 2^4 for Q3.4) means **less precision** but **larger range**.

Example Comparisons:

Q-Format	Scaling Factor	Smallest Representable Step
Q1.15	$2^{15} = 32768$	$\frac{1}{32768} = 0.00003$
Q3.12	$2^{12} = 4096$	$\frac{1}{4096} = 0.00024$
Q5.10	$2^{10} = 1024$	$\frac{1}{1024} = 0.00098$

To represent -0.0265 in Q1.15 format, we follow these steps:

1. **Convert the positive value of 0.0265** into Q1.15 format (as done previously):
 - $0.0265 \times 32768 = 866.56$
 - Rounding gives 867, and in binary: $0000\ 0011\ 1000\ 0011_2$.
2. **Represent the negative value** by taking the 2's complement of the positive value's binary form:
 - The binary of 867 is $0000\ 0011\ 1000\ 0011_2$.
 - To find the 2's complement:
 - Invert the bits: $1111\ 1100\ 0111\ 1100_2$
 - Add 1: $1111\ 1100\ 0111\ 1101_2$

So, the 2's complement of 867 (for -0.0265) is $1111\ 1100\ 0111\ 1101_2$.

3. **Convert the binary form to hexadecimal:**
 - $1111\ 1100\ 0111\ 1101_2$ is equivalent to **0xFC7D** in hexadecimal.

Thus, -0.0265 in Q1.15 format in hexadecimal is **0xFC7D**.